

Java binary serialization review

Sirenko A.

December 2, 2014

2014

Requirements

We use serialization in 2 cases:

- 1 to store;
- 2 to transfer;

Main features to consider:

- 1 serialize/deserialize speed;
- 2 serialized size;
- 3 stability to change;
- 4 laboriousness;
- 5 platform-dependence;

JDK tools

- Standard serialization: smallest effort / worst performance;
- Externalizable serialization: full control over serialization, define your binary format, no features;

| serialization | speed | size | stability to change | laboriousness | platform dependence |
|----------------|-------|-------|------------------------|---------------|------------------------|
| standard | poor | big | average | low | only java |
| externalizable | fast | small | poor | high | independent |

Table : Standard serialization vs Externalizable

Why?

Motivation:

- conveniently simple and fast approaches;
- better adopted for our data: optional params, variable-length numbers, aliases;
- platform-independent: same description of our data for different platforms.

Here we consider:

- protobuf - <https://code.google.com/p/protobuf/>;
- avro - <http://avro.apache.org/>;
- kryo - <https://github.com/EsotericSoftware/kryo>.

Protobuf

Protobuf info

Avro

Avro info

Kryo

Kryo info

Comparison

Comparison placeholder

Kryo vs. Std in queue